

_sentiment_analysis_assignment

July 10, 2023

```
[ ]: !pip install transformers
```

Collecting transformers

Downloading transformers-4.11.3-py3-none-any.whl (2.9 MB)

| 2.9 MB 4.1 MB/s

Requirement already satisfied: tqdm>=4.27 in

/usr/local/lib/python3.7/dist-packages (from transformers) (4.62.3)

Requirement already satisfied: importlib-metadata in

/usr/local/lib/python3.7/dist-packages (from transformers) (4.8.1)

Collecting huggingface-hub>=0.0.17

Downloading huggingface_hub-0.0.19-py3-none-any.whl (56 kB)

| 56 kB 5.0 MB/s

Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from transformers) (2.23.0)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.7/dist-packages (from transformers) (21.0)

Collecting tokenizers<0.11,>=0.10.1

Downloading tokenizers-0.10.3-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_12_x86_64.manylinux2010_x86_64.whl (3.3 MB)

| 3.3 MB 35.4 MB/s

Requirement already satisfied: regex!=2019.12.17 in

/usr/local/lib/python3.7/dist-packages (from transformers) (2019.12.20)

Collecting pyyaml>=5.1

Downloading PyYAML-6.0-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_12_x86_64.manylinux2010_x86_64.whl (596 kB)

| 596 kB 41.6 MB/s

Requirement already satisfied: numpy>=1.17 in

/usr/local/lib/python3.7/dist-packages (from transformers) (1.19.5)

Collecting sacremoses

Downloading sacremoses-0.0.46-py3-none-any.whl (895 kB)

| 895 kB 48.7 MB/s

Requirement already satisfied: filelock in /usr/local/lib/python3.7/dist-packages (from transformers) (3.3.0)

Requirement already satisfied: typing-extensions in

/usr/local/lib/python3.7/dist-packages (from huggingface-hub>=0.0.17->transformers) (3.7.4.3)

Requirement already satisfied: pyparsing>=2.0.2 in

/usr/local/lib/python3.7/dist-packages (from packaging>=20.0->transformers)

(2.4.7)

Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from importlib-metadata->transformers) (3.6.0)

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests->transformers) (2.10)

Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests->transformers) (3.0.4)

Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests->transformers) (1.24.3)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests->transformers) (2021.5.30)

Requirement already satisfied: joblib in /usr/local/lib/python3.7/dist-packages (from sacremoses->transformers) (1.0.1)

Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from sacremoses->transformers) (1.15.0)

Requirement already satisfied: click in /usr/local/lib/python3.7/dist-packages (from sacremoses->transformers) (7.1.2)

Installing collected packages: pyyaml, tokenizers, sacremoses, huggingface-hub, transformers

Attempting uninstall: pyyaml

Found existing installation: PyYAML 3.13

Uninstalling PyYAML-3.13:

Successfully uninstalled PyYAML-3.13

Successfully installed huggingface-hub-0.0.19 pyyaml-6.0 sacremoses-0.0.46 tokenizers-0.10.3 transformers-4.11.3

```
[ ]: import transformers
from transformers import BertModel, BertTokenizer, AdamW, \
    get_linear_schedule_with_warmup
import torch

import numpy as np
import pandas as pd
import seaborn as sns
from pylab import rcParams
import matplotlib.pyplot as plt
from matplotlib import rc
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from collections import defaultdict
from textwrap import wrap

from torch import nn, optim
from torch.utils.data import Dataset, DataLoader
import torch.nn.functional as F

import warnings
```

```
warnings.filterwarnings('ignore')
```

```
[ ]: %matplotlib inline
      %config InlineBackend.figure_format='retina'
```

```
[ ]: sns.set(style='whitegrid', palette='muted', font_scale=1.2)
      HAPPY_COLORS_PALETTE = ["#01BEFE", "#FFDD00", "#FF7D00", "#FF006D", "#ADFF02",
                              ↵ "#8F00FF"]
      sns.set_palette(sns.color_palette(HAPPY_COLORS_PALETTE))
      rcParams['figure.figsize'] = 12, 8
      RANDOM_SEED = 42
      np.random.seed(RANDOM_SEED)
      torch.manual_seed(RANDOM_SEED)
      # device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")

      # If there's a GPU available...
      if torch.cuda.is_available():

          # Tell PyTorch to use the GPU.
          device = torch.device("cuda")

          print('There are %d GPU(s) available.' % torch.cuda.device_count())

          print('We will use the GPU:', torch.cuda.get_device_name(0))

      # If not...
      else:
          print('No GPU available, using the CPU instead.')
          device = torch.device("cpu")
```

There are 1 GPU(s) available.

We will use the GPU: Tesla K80

```
[ ]: !gdown --id 1S6qMioqPJjyBLpLVz4gmRTnJHnjitnuV
      !gdown --id 1zdmewp7ayS4js4VtrJEHzAheSW-5NBZv
```

Downloading...

From: <https://drive.google.com/uc?id=1S6qMioqPJjyBLpLVz4gmRTnJHnjitnuV>

To: /content/apps.csv

100% 134k/134k [00:00<00:00, 46.2MB/s]

Downloading...

From: <https://drive.google.com/uc?id=1zdmewp7ayS4js4VtrJEHzAheSW-5NBZv>

To: /content/reviews.csv

100% 7.17M/7.17M [00:00<00:00, 22.9MB/s]

```
[ ]: df = pd.read_csv("reviews.csv")
df.head()
```

```
[ ]:      userName  ...  appId
0  Andrew Thomas  ...  com.anydo
1   Craig Haines  ...  com.anydo
2  steven adkins  ...  com.anydo
3  Lars Panzerbjørn  ...  com.anydo
4   Scott Prewitt  ...  com.anydo
```

[5 rows x 11 columns]


```
[ ]: df.shape
```

```
[ ]: (15746, 11)
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15746 entries, 0 to 15745
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   userName              15746 non-null  object
1   userImage             15746 non-null  object
2   content               15746 non-null  object
3   score                 15746 non-null  int64
4   thumbsUpCount         15746 non-null  int64
5   reviewCreatedVersion  13533 non-null  object
6   at                    15746 non-null  object
7   replyContent          7367 non-null   object
8   repliedAt             7367 non-null   object
9   sortOrder             15746 non-null  object
10  appId                 15746 non-null  object
dtypes: int64(2), object(9)
memory usage: 1.3+ MB
```


```
[ ]: sns.countplot(df.score)
plt.xlabel('review score');
```



```
[ ]: def to_sentiment(rating):  
    rating = int(rating)  
    if rating <= 2:  
        return 0  
    elif rating == 3:  
        return 1  
    else:  
        return 2  
  
df['sentiment'] = df.score.apply(to_sentiment)
```

```
[ ]: class_names = ['negative', 'neutral', 'positive']
```

```
[ ]: ax = sns.countplot(df.sentiment)  
plt.xlabel('review sentiment')  
ax.set_xticklabels(class_names);
```



```
[ ]: PRE_TRAINED_MODEL_NAME = 'bert-base-cased'
```

```
[ ]: tokenizer = BertTokenizer.from_pretrained(PRE_TRAINED_MODEL_NAME)
```

```
Downloading: 0%|          | 0.00/208k [00:00<?, ?B/s]
```

```
Downloading: 0%|          | 0.00/29.0 [00:00<?, ?B/s]
```

```
Downloading: 0%|          | 0.00/426k [00:00<?, ?B/s]
```

```
Downloading: 0%|          | 0.00/570 [00:00<?, ?B/s]
```

```
[ ]: sample_txt = 'When was I last outside? I am stuck at home for 2 weeks.'
```

```
[ ]: tokens = tokenizer.tokenize(sample_txt)
token_ids = tokenizer.convert_tokens_to_ids(tokens)
```

```
print(f' Sentence: {sample_txt}')
```

```
print(f'  Tokens: {tokens}')
```

```
print(f'Token IDs: {token_ids}')
```

```
Sentence: When was I last outside? I am stuck at home for 2 weeks.
```

```
Tokens: ['When', 'was', 'I', 'last', 'outside', '?', 'I', 'am', 'stuck',
'at', 'home', 'for', '2', 'weeks', '.']
```

```
Token IDs: [1332, 1108, 146, 1314, 1796, 136, 146, 1821, 5342, 1120, 1313, 1111,
123, 2277, 119]
```

```
[ ]: tokenizer.sep_token, tokenizer.sep_token_id
```

```
[ ]: ('[SEP]', 102)
```

```
[ ]: tokenizer.cls_token, tokenizer.cls_token_id
```

```
[ ]: ('[CLS]', 101)
```

```
[ ]: tokenizer.pad_token, tokenizer.pad_token_id
```

```
[ ]: ('[PAD]', 0)
```

```
[ ]: tokenizer.unk_token, tokenizer.unk_token_id
```

```
[ ]: ('[UNK]', 100)
```

```
[ ]: encoding = tokenizer.encode_plus(  
    sample_txt,  
    max_length=32,  
    add_special_tokens=True, # Add '[CLS]' and '[SEP]'  
    return_token_type_ids=False,  
    pad_to_max_length=True,  
    return_attention_mask=True,  
    return_tensors='pt', # Return PyTorch tensors  
)  
  
encoding.keys()
```

Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate examples to max length. Defaulting to 'longest_first' truncation strategy. If you encode pairs of sequences (GLUE-style) with the tokenizer you can select this strategy more precisely by providing a specific strategy to `truncation`.

```
[ ]: dict_keys(['input_ids', 'attention_mask'])
```

```
[ ]: print(len(encoding['input_ids'][0]))  
encoding['input_ids'][0]
```

32

```
[ ]: tensor([ 101, 1332, 1108,  146, 1314, 1796,  136,  146, 1821, 5342, 1120, 1313,  
            1111,  123, 2277,  119,  102,   0,   0,   0,   0,   0,   0,   0,  
             0,   0,   0,   0,   0,   0,   0,   0])
```

```
[ ]: print(len(encoding['attention_mask'][0]))  
encoding['attention_mask']
```

```
[ ]: tensor([[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0]])
```

```
[ ]: tokenizer.convert_ids_to_tokens(encoding['input_ids'][0])
```

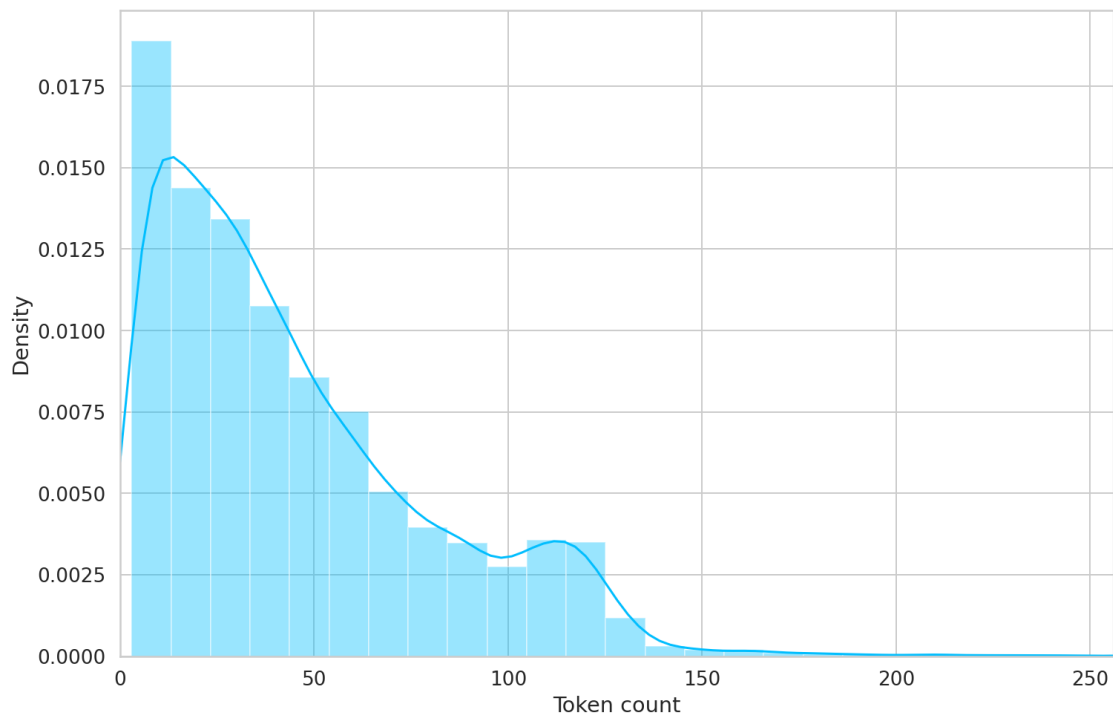
[illegible]

```
[ ]: token_lens = []

for content in df.content:
    tokens = tokenizer.encode(content, max_length=512)
    token_lens.append(len(tokens))
```



```
[ ]: sns.distplot(token_lens)
plt.xlim([0, 256]);
plt.xlabel('Token count');
```



```
[ ]: MAX_LEN = 160
```

```
[ ]: class GPReviewDataset(Dataset):
    def __init__(self, reviews, targets, tokenizer, max_len):
        self.reviews = reviews
        self.targets = targets
        self.tokenizer = tokenizer
        self.max_len = max_len

    def __len__(self):
        return len(self.reviews)

    def __getitem__(self, item):
        review = str(self.reviews[item])
        target = self.targets[item]

        encoding = self.tokenizer.encode_plus(
            review,
            add_special_tokens=True,
            max_length=self.max_len,
```

```

        return_token_type_ids=False,
        pad_to_max_length=True,
        return_attention_mask=True,
        return_tensors='pt',
    )

    return {
        'review_text': review,
        'input_ids': encoding['input_ids'].flatten(),
        'attention_mask': encoding['attention_mask'].flatten(),
        'targets': torch.tensor(target, dtype=torch.long)
    }

```

```

[ ]: df_train, df_test = train_test_split(df, test_size=0.3,
    ↪random_state=RANDOM_SEED)

df_val, df_test = train_test_split(df_test, test_size=0.5,
    ↪random_state=RANDOM_SEED)

```

```
[ ]: df_train.shape
```

```
[ ]: (11022, 12)
```

```
[ ]: df_val.shape
```

```
[ ]: (2362, 12)
```

```
[ ]: df_test.shape
```

```
[ ]: (2362, 12)
```

```

[ ]: def create_data_loader(df, tokenizer, max_len, batch_size):
    ds = GPReviewDataset(
        reviews=df.content.to_numpy(),
        targets=df.sentiment.to_numpy(),
        tokenizer=tokenizer,
        max_len=max_len
    )

    return DataLoader(
        ds,
        batch_size=batch_size,
        num_workers=2
    )

```

```
[ ]: BATCH_SIZE = 16
```

```
train_data_loader = create_data_loader(df_train, tokenizer, MAX_LEN, BATCH_SIZE)
val_data_loader = create_data_loader(df_val, tokenizer, MAX_LEN, BATCH_SIZE)
test_data_loader = create_data_loader(df_test, tokenizer, MAX_LEN, BATCH_SIZE)
```

```
[ ]: data = next(iter(train_data_loader))
data.keys()
```

```
[ ]: dict_keys(['review_text', 'input_ids', 'attention_mask', 'targets'])
```

```
[ ]: print(data['input_ids'].shape)
print(data['attention_mask'].shape)
print(data['targets'].shape)
```

```
torch.Size([16, 160])
torch.Size([16, 160])
torch.Size([16])
```

```
[ ]: bert_model = BertModel.from_pretrained(PRE_TRAINED_MODEL_NAME)
```

```
Downloading: 0%|          | 0.00/416M [00:00<?, ?B/s]
```

Some weights of the model checkpoint at bert-base-cased were not used when initializing BertModel: ['cls.predictions.transform.LayerNorm.weight', 'cls.seq_relationship.bias', 'cls.seq_relationship.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight', 'cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.bias']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

```
[ ]: last_hidden_state, pooled_output = bert_model(
    input_ids=encoding['input_ids'], attention_mask=encoding['attention_mask'],
    return_dict = False)
```

```
[ ]: # last_hidden_state=dict['last_hidden_state']
# pooled_output=dict['pooler_output']

last_hidden_state.shape
```

```
[ ]: torch.Size([1, 32, 768])
```

```
[ ]: bert_model.config.hidden_size
```

```
[ ]: 768
```

```
[ ]: pooled_output.shape
```

```
[ ]: torch.Size([1, 768])
```

```
[ ]: class SentimentClassifier(nn.Module):
    def __init__(self, n_classes):
        super(SentimentClassifier, self).__init__()
        self.bert = BertModel.from_pretrained(PRE_TRAINED_MODEL_NAME)
        self.drop = nn.Dropout(p=0.3)
        self.out = nn.Linear(self.bert.config.hidden_size, n_classes)

    # def forward(self, input_ids, attention_mask):
    #     _, pooled_output = self.bert(input_ids=input_ids,
    # ↪attention_mask=attention_mask)
    #     output = self.drop(pooled_output)

    #     return self.out(output)

    def forward(self, input_ids, attention_mask):
        returned = self.bert(input_ids=input_ids, attention_mask=attention_mask)

        pooled_output = returned["pooler_output"]
        output = self.drop(pooled_output)

        return self.out(output)
```

```
[ ]: model = SentimentClassifier(len(class_names))
model = model.to(device)
```

Some weights of the model checkpoint at bert-base-cased were not used when initializing BertModel: ['cls.predictions.transform.LayerNorm.weight', 'cls.seq_relationship.bias', 'cls.seq_relationship.weight', 'cls.predictions.transform.dense.bias', 'cls.predictions.decoder.weight', 'cls.predictions.bias', 'cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.bias']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

```
[ ]: input_ids = data['input_ids'].to(device)
attention_mask = data['attention_mask'].to(device)
print(input_ids.shape) # batch size x seq length
print(attention_mask.shape) # batch size x seq length
```

```
torch.Size([16, 160])
torch.Size([16, 160])
```

```
[ ]: F.softmax(model(input_ids, attention_mask), dim=1)
```

```
[ ]: tensor([[0.2236, 0.4834, 0.2931],
            [0.2244, 0.3195, 0.4561],
            [0.3162, 0.2818, 0.4020],
            [0.2029, 0.4641, 0.3330],
            [0.5511, 0.2338, 0.2151],
            [0.2239, 0.4568, 0.3194],
            [0.2738, 0.3206, 0.4055],
            [0.4069, 0.1912, 0.4020],
            [0.4032, 0.1856, 0.4113],
            [0.3382, 0.1997, 0.4621],
            [0.3512, 0.2262, 0.4226],
            [0.3372, 0.2153, 0.4474],
            [0.2272, 0.3774, 0.3954],
            [0.2966, 0.3315, 0.3719],
            [0.3520, 0.2277, 0.4203],
            [0.1511, 0.3517, 0.4973]], device='cuda:0', grad_fn=<SoftmaxBackward>)
```

```
[ ]: EPOCHS = 10
optimizer = AdamW(model.parameters(), lr=2e-5, correct_bias=False)
total_steps = len(train_data_loader) * EPOCHS
scheduler = get_linear_schedule_with_warmup(
    optimizer,
    num_warmup_steps=0,
    num_training_steps=total_steps
)
loss_fn = nn.CrossEntropyLoss().to(device)
```

```
[ ]: def train_epoch(model, data_loader, loss_fn, optimizer, device, scheduler,
    ↪n_examples):
    model = model.train()

    losses = []
    correct_predictions = 0

    for d in data_loader:
        input_ids = d["input_ids"].to(device)
        attention_mask = d["attention_mask"].to(device)
        targets = d["targets"].to(device)

        outputs = model(
            input_ids=input_ids,
            attention_mask=attention_mask
```

```

    )

    _, preds = torch.max(outputs, dim=1)
    loss = loss_fn(outputs, targets)

    correct_predictions += torch.sum(preds == targets)
    losses.append(loss.item())

    loss.backward()
    nn.utils.clip_grad_norm_(model.parameters(), max_norm=1.0)
    optimizer.step()
    scheduler.step()
    optimizer.zero_grad()

    return correct_predictions.double() / n_examples, np.mean(losses)

```

```

[ ]: def eval_model(model, data_loader, loss_fn, device, n_examples):
    model = model.eval()

    losses = []
    correct_predictions = 0

    with torch.no_grad():
        for d in data_loader:
            input_ids = d["input_ids"].to(device)
            attention_mask = d["attention_mask"].to(device)
            targets = d["targets"].to(device)

            outputs = model(
                input_ids=input_ids,
                attention_mask=attention_mask
            )
            _, preds = torch.max(outputs, dim=1)

            loss = loss_fn(outputs, targets)

            correct_predictions += torch.sum(preds == targets)
            losses.append(loss.item())

    return correct_predictions.double() / n_examples, np.mean(losses)

```

```

[ ]: from google.colab import drive
    drive.mount('/content/drive')

```

Mounted at /content/drive

```

[ ]: %%time

history = defaultdict(list)
best_accuracy = 0

for epoch in range(EPOCHS):
    print(f'Epoch {epoch + 1}/{EPOCHS}')
    print('-' * 10)

    train_acc, train_loss = train_epoch(
        model,
        train_data_loader,
        loss_fn,
        optimizer,
        device,
        scheduler,
        len(df_train)
    )

    print(f'Train loss {train_loss} --- Train accuracy {train_acc}')

    val_acc, val_loss = eval_model(
        model,
        val_data_loader,
        loss_fn,
        device,
        len(df_val)
    )

    print(f'Val loss {val_loss} --- Val accuracy {val_acc}')
    print()

    history['train_acc'].append(train_acc)
    history['train_loss'].append(train_loss)
    history['val_acc'].append(val_acc)
    history['val_loss'].append(val_loss)

    if val_acc > best_accuracy:
        model_save_name = 'best_model_state.pt'
        path = F"/content/drive/My Drive/{model_save_name}"
        torch.save(model.state_dict(), path)

        # torch.save(model.state_dict(), 'best_model_state.bin')
        best_accuracy = val_acc

```

Epoch 1/10

Train loss 0.7674604126607731 --- Train accuracy 0.6470694973688985
Val loss 0.6753895640171863 --- Val accuracy 0.7146486028789162

Epoch 2/10

Train loss 0.4834810435577297 --- Train accuracy 0.8122845218653602
Val loss 0.67174454239776 --- Val accuracy 0.7641828958509738

Epoch 3/10

Train loss 0.28917152965806303 --- Train accuracy 0.9043730720377426
Val loss 0.8130055506638175 --- Val accuracy 0.7895850973751058

Epoch 4/10

Train loss 0.18364281476041402 --- Train accuracy 0.9466521502449646
Val loss 0.9616580203354888 --- Val accuracy 0.8044030482641829

Epoch 5/10

Train loss 0.13613030183980923 --- Train accuracy 0.9620758483033932
Val loss 0.9913723825082196 --- Val accuracy 0.8243014394580863

Epoch 6/10

Train loss 0.09584020892865598 --- Train accuracy 0.9731446198512067
Val loss 1.0787120965762518 --- Val accuracy 0.8247248094834886

Epoch 7/10

Train loss 0.07124643715384724 --- Train accuracy 0.9781346398112866
Val loss 1.1752637787159639 --- Val accuracy 0.8268416596104995

Epoch 8/10

Train loss 0.061142171285499064 --- Train accuracy 0.982308111050626
Val loss 1.2530784735851763 --- Val accuracy 0.8272650296359018

Epoch 9/10

Train loss 0.04363285115154944 --- Train accuracy 0.9869352204681546
Val loss 1.278646860217221 --- Val accuracy 0.8302286198137172


Epoch 10/10

Train loss 0.04183529664257373 --- Train accuracy 0.9872981310107058
Val loss 1.2865800096325395 --- Val accuracy 0.8293818797629128

CPU times: user 1h 49min 22s, sys: 1min, total: 1h 50min 23s
Wall time: 1h 50min 40s

```
[ ]: plt.plot(history['train_acc'], label='train accuracy')
plt.plot(history['val_acc'], label='validation accuracy')

plt.title('Training history')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.ylim([0, 1]);
```



```
[ ]: test_acc, _ = eval_model(
    model,
    test_data_loader,
    loss_fn,
    device,
    len(df_test)
)

test_acc.item()
```

```
[ ]: 0.8209144792548687
```

```
[ ]: def get_predictions(model, data_loader):
    model = model.eval()

    review_texts = []
    predictions = []
    prediction_probs = []
    real_values = []

    with torch.no_grad():
        for d in data_loader:
            texts = d["review_text"]
            input_ids = d["input_ids"].to(device)
            attention_mask = d["attention_mask"].to(device)
            targets = d["targets"].to(device)

            outputs = model(
                input_ids=input_ids,
                attention_mask=attention_mask
            )
            _, preds = torch.max(outputs, dim=1)

            probs = F.softmax(outputs, dim=1)

            review_texts.extend(texts)
            predictions.extend(preds)
            prediction_probs.extend(probs)
            real_values.extend(targets)

    predictions = torch.stack(predictions).cpu()
    prediction_probs = torch.stack(prediction_probs).cpu()
    real_values = torch.stack(real_values).cpu()
    return review_texts, predictions, prediction_probs, real_values
```

```
[ ]: y_review_texts, y_pred, y_pred_probs, y_test = get_predictions(
    model,
    test_data_loader
)
```


```
[ ]: print(classification_report(y_test, y_pred, target_names=class_names))
```

	precision	recall	f1-score	support
negative	0.86	0.82	0.84	767
neutral	0.73	0.77	0.75	731
positive	0.86	0.87	0.87	864
accuracy			0.82	2362

macro avg	0.82	0.82	0.82	2362
weighted avg	0.82	0.82	0.82	2362

```
[ ]: def show_confusion_matrix(confusion_matrix):
    hmap = sns.heatmap(confusion_matrix, annot=True, fmt="d", cmap="Blues")
    hmap.yaxis.set_ticklabels(hmap.yaxis.get_ticklabels(), rotation=0,
    ↪ha='right')
    hmap.xaxis.set_ticklabels(hmap.xaxis.get_ticklabels(), rotation=30,
    ↪ha='right')
    plt.ylabel('True sentiment')
    plt.xlabel('Predicted sentiment');

cm = confusion_matrix(y_test, y_pred)
df_cm = pd.DataFrame(cm, index=class_names, columns=class_names)
show_confusion_matrix(df_cm)
```



```
[ ]: idx = 2

review_text = y_review_texts[idx]
true_sentiment = y_test[idx]
pred_df = pd.DataFrame({
```

```

    'class_names': class_names,
    'values': y_pred_probs[idx]
})

```

```

[ ]: print("\n".join(wrap(review_text)))
      print()
      print(f'True sentiment: {class_names[true_sentiment]}')

```


I have tied a few calenders. This one is good for multiple accounts.
Easy to turn them off and on. Its quite clear on a phone screen

True sentiment: positive

```

[ ]: sns.barplot(x='values', y='class_names', data=pred_df, orient='h')
      plt.ylabel('sentiment')
      plt.xlabel('probability')
      plt.xlim([0, 1]);

```



```

[ ]: review_text = "maybe i like you."

```

```

[ ]: encoded_review = tokenizer.encode_plus(
      review_text,
      max_length=MAX_LEN,
      add_special_tokens=True,

```

```

    return_token_type_ids=False,
    pad_to_max_length=True,
    return_attention_mask=True,
    return_tensors='pt',
)

```

```

[ ]: input_ids = encoded_review['input_ids'].to(device)
    attention_mask = encoded_review['attention_mask'].to(device)

    output = model(input_ids, attention_mask)
    _, prediction = torch.max(output, dim=1)

    print(f'Review text: {review_text}')
    print(f'Sentiment : {class_names[prediction]}')

```

Review text: maybe i like you.
Sentiment : neutral

```

[ ]: model_save_name = 'best_model_state.bin'
    path = F"/content/drive/My Drive/{model_save_name}"
    torch.save(model.state_dict(), path)
    # model.load_state_dict(torch.load(path))

```

```

[ ]: model_save_name = 'best_model_state.bin'
    path = F"/content/drive/My Drive/{model_save_name}"
    model.load_state_dict(torch.load(path))

```

[]: <All keys matched successfully>

```

[ ]: sentiment_test = "maybe i like you."

```

```

[ ]: encoded_review = tokenizer.encode_plus(
    sentiment_test,
    max_length=MAX_LEN,
    add_special_tokens=True,
    return_token_type_ids=False,
    pad_to_max_length=True,
    return_attention_mask=True,
    return_tensors='pt',
)

```

```

[ ]: input_ids = encoded_review['input_ids'].to(device)
    attention_mask = encoded_review['attention_mask'].to(device)

    output = model(input_ids, attention_mask)
    _, prediction = torch.max(output, dim=1)

```

```
print(f'Review text: {sentiment_test}')  
print(f'Sentiment : {class_names[prediction]}')
```

```
Review text: maybe i like you.  
Sentiment : neutral
```

```
[ ]:
```